

# Branch Prediction in ARM Processors

## Contents

<b>1 Purpose of Branch Predictor</b>	<b>2</b>
<b>2 Types of Branch Prediction in ARM Processors</b>	<b>2</b>
2.1 Static Branch Prediction . . . . .	2
2.2 Dynamic Branch Prediction . . . . .	3
<b>3 Branch Prediction Architecture in ARM Processors</b>	<b>3</b>
3.1 Branch Target Buffer (BTB) . . . . .	3
3.2 Pattern History Table (PHT) . . . . .	3
3.3 Global and Local Predictors . . . . .	4
3.4 Return Address Stack (RAS) . . . . .	4
<b>4 How Branch Prediction Works in ARM</b>	<b>4</b>
4.1 Pipeline Stages with Branch Prediction . . . . .	4
<b>5 ARMv8 and ARMv9 Enhancements</b>	<b>5</b>
5.1 ARMv8-A (AArch64) . . . . .	5
5.2 ARMv9 . . . . .	5
<b>6 Key Features in ARM Branch Predictors</b>	<b>5</b>
6.1 Low Latency Prediction . . . . .	5
6.2 Energy Efficiency . . . . .	5
6.3 High Accuracy . . . . .	5
<b>7 Common Metrics for Evaluating Branch Predictors</b>	<b>5</b>
<b>8 Example of Branch Prediction in ARM</b>	<b>6</b>
8.1 Backward Branch in a Loop . . . . .	6
<b>9 Conclusion</b>	<b>6</b>

# Introduction

Modern processors use **branch prediction** to minimize pipeline stalls caused by conditional instructions (e.g., loops, if-else). In ARM processors, which power mobile and embedded systems, branch prediction is a key feature for achieving **high performance**, **energy efficiency**, and **accuracy**.

## The Need for Branch Prediction

- **Problem:** When a branch instruction is encountered, the processor must decide:
  - **Taken:** Execute the branch.
  - **Not-taken:** Continue executing sequential instructions.
- **Challenge:** Deciding which path to take delays execution.
- **Solution:** Use branch prediction to predict the outcome and fetch instructions speculatively.

## 1 Purpose of Branch Predictor

The purpose of a branch predictor is to:

1. Reduce pipeline stalls caused by branch mispredictions.
2. Maintain high throughput in pipelined architectures.
3. Optimize performance in **mobile**, **IoT**, and **high-performance** systems.

## Pipeline Delay

Without branch prediction, the processor stalls to determine the outcome of a branch:

$$\text{Delay (cycles)} = \text{Pipeline Depth} \times \text{Branch Misprediction Rate}$$

## 2 Types of Branch Prediction in ARM Processors

ARM processors implement two main types of branch prediction: **static** and **dynamic**.

### 2.1 Static Branch Prediction

**Static Prediction Rules:**

- **Backward Branches:** Predicted as **taken** (e.g., loops).
- **Forward Branches:** Predicted as **not-taken**.

**Advantages:**

- Simple to implement.

- No hardware overhead.

**Limitations:**

- Cannot adapt to branch behavior.

## 2.2 Dynamic Branch Prediction

Dynamic prediction adapts to **runtime branch behavior** using:

- **History-based predictors.**
- **Hybrid models** (e.g., global + local history).

## 3 Branch Prediction Architecture in ARM Processors

ARM processors utilize **multi-level architectures** for branch prediction.

### 3.1 Branch Target Buffer (BTB)

The **BTB** is a cache storing:

- Branch instruction addresses.
- Predicted target addresses.

When a branch instruction is fetched:

$$\text{Predicted Target Address} = \text{BTB}[\text{Branch PC}]$$

**Role of BTB:**

- Reduces latency by prefetching target instructions.
- Updates dynamically with new branches.

### 3.2 Pattern History Table (PHT)

The **PHT** tracks branch outcomes using **2-bit saturating counters**:

1. **Strongly Taken (ST)**: Always predict taken.
2. **Weakly Taken (WT)**: Predict taken, but can switch.
3. **Weakly Not-Taken (WNT)**: Predict not-taken, but can switch.
4. **Strongly Not-Taken (SNT)**: Always predict not-taken.

$$\text{Next State} = f(\text{Current State}, \text{Branch Outcome})$$

Example Transition:

$$\text{WT} \xrightarrow{\text{Not-Taken}} \text{WNT}$$

### 3.3 Global and Local Predictors

- **Global History Register (GHR):**

- Tracks outcomes of the last  $N$  branches.
- Correlates branch predictions with global patterns.

Example:

$$\text{GHR} = [T, T, NT, T] \quad (\text{T: Taken, NT: Not-Taken})$$

- **Local History Table (LHT):**

- Tracks branch-specific history.
- Useful for branches with unique behavior.

### 3.4 Return Address Stack (RAS)

The **RAS** predicts the return address for function calls:

1. Pushes return addresses during function calls.
2. Pops addresses during returns.

## 4 How Branch Prediction Works in ARM

### 4.1 Pipeline Stages with Branch Prediction

1. **Fetch Stage:**

- The branch predictor identifies branch instructions.
- Checks the **BTB** and **PHT** for prediction data.

2. **Prediction Stage:**

- If a branch is found in the BTB:

Target Address  $\rightarrow$  Prefetch

- Otherwise:
  - Use static prediction.

3. **Execution Stage:**

- Executes the branch.
- If the prediction was wrong:
  - Flush the pipeline.
  - Fetch correct instructions.

## 5 ARMv8 and ARMv9 Enhancements

### 5.1 ARMv8-A (AArch64)

Key Features:

- **Improved BTB:** Larger size to store more branches.
- **Advanced Predictors:** Combines local and global history.

### 5.2 ARMv9

New Enhancements:

- **Neural Predictors:** Leverages machine learning for irregular patterns.
- **Thread-Specific Predictors:** Separate tables for multi-threaded execution.

## 6 Key Features in ARM Branch Predictors

### 6.1 Low Latency Prediction

Predictions occur at the fetch stage to minimize stalls.

### 6.2 Energy Efficiency

Optimized for mobile devices where power consumption is critical.

### 6.3 High Accuracy

Combines global and local predictors to improve prediction rates.

## 7 Common Metrics for Evaluating Branch Predictors

1. **Branch Prediction Accuracy (BPA):**

$$\text{BPA} = \frac{\text{Correct Predictions}}{\text{Total Branches}} \times 100$$

2. **Misprediction Penalty:** Cycles lost due to incorrect predictions.

3. **BTB Hit Rate:**

$$\text{BTB Hit Rate} = \frac{\text{Branches Found in BTB}}{\text{Total Branches}}$$

## 8 Example of Branch Prediction in ARM

### 8.1 Backward Branch in a Loop

```
MOV R0, #0           ; Initialize counter
LOOP:
ADD R0, R0, #1       ; Increment counter
CMP R0, #10          ; Compare counter with 10
BNE LOOP             ; Branch if not equal (backward branch)
```

- **Static Prediction:** Predicts as **taken**.
- **Dynamic Prediction:** Adapts to runtime behavior for higher accuracy.

## 9 Conclusion

ARM processors use advanced branch prediction mechanisms to achieve high performance, low latency, and energy efficiency. Techniques like **dynamic prediction**, **neural predictors**, and **hybrid models** ensure ARM architectures excel in modern workloads.



# Speculative Execution in ARM Processors

## Contents

	3
1 Purpose of Speculative Execution	3
2 How Speculative Execution Works	3
3 Mathematical Representation of Speculative Execution	3
3.1 Pipeline Performance Without Speculation	4
3.2 Pipeline Performance With Speculation	4
3.3 Performance Gain	4
4 Components Supporting Speculative Execution in ARM Processors	4
4.1 Branch Predictors	4
4.2 Branch Target Buffer (BTB)	4
4.3 Reorder Buffer (ROB)	4
4.4 Return Address Stack (RAS)	4
5 ARM Speculative Execution Mechanism	5
5.1 Step-by-Step Process	5
6 Equations for Speculative Execution in ARM	5
6.1 Correct Prediction	5
6.2 Incorrect Prediction	5
6.3 Overall Execution Time	5
7 Security Implications of Speculative Execution	5
7.1 Side-Channel Attacks	6
7.2 Mitigations in ARM	6
8 ARM Architectures with Speculative Execution	6
8.1 ARMv7	6
8.2 ARMv8	6
8.3 ARMv9	6
9 Advantages of Speculative Execution	6
10 Challenges of Speculative Execution	6
11 Example: Speculative Execution in ARM Assembly	7

<b>12 Metrics for Evaluating Speculative Execution</b>	<b>7</b>
12.1 Speculation Efficiency (SE) . . . . .	7
12.2 Branch Prediction Accuracy (BPA) . . . . .	7
12.3 Misprediction Penalty (MP) . . . . .	7
<b>13 Future of Speculative Execution in ARM</b>	<b>7</b>
<b>14 Conclusion</b>	<b>8</b>



# Introduction

**Speculative Execution** is a performance optimization technique used in modern processors, including ARM processors, to execute instructions before the final outcome of a decision (e.g., a branch) is known. By predicting the path of execution, the processor can execute instructions speculatively, reducing delays caused by pipeline stalls.

## 1 Purpose of Speculative Execution

### 1. Increase Instruction Throughput:

- Minimizes idle time in the pipeline.
- Keeps the pipeline filled with instructions.

### 2. Reduce Branch Penalty:

- Branch instructions introduce delays due to uncertainty about which path to execute.
- Speculative execution helps mitigate this delay.

### 3. Optimize Multi-Stage Pipelines:

- ARM processors with deeper pipelines benefit significantly from speculation as it minimizes stalls.

## 2 How Speculative Execution Works

### 1. Instruction Fetch Stage:

- The processor fetches instructions for both possible outcomes of a branch.

### 2. Prediction Stage:

- The **branch predictor** predicts the most likely branch to be taken.
- Example: Global History Register (GHR) or Branch Target Buffer (BTB).

### 3. Execution Stage:

- Executes instructions along the predicted path speculatively.
- If the prediction is correct, the results are committed.
- If incorrect, the pipeline is flushed, and the correct path is executed.

## 3 Mathematical Representation of Speculative Execution

### 3.1 Pipeline Performance Without Speculation

Let

- $I$ : Total number of instructions.
- $C$ : Cycles per instruction (CPI).
- $P_m$ : Probability of branch misprediction.
- $T_m$ : Misprediction penalty in cycles.

The total execution time  $T_{\text{no-speculation}}$  is given by:

$$T_{\text{no-speculation}} = I \times C$$

### 3.2 Pipeline Performance With Speculation

Speculative execution introduces potential branch misprediction penalties. The execution time becomes:

$$T_{\text{speculation}} = I \times C + P_m \times T_m \times \text{Branch Instructions}$$

### 3.3 Performance Gain

$$\text{Performance Gain} = \frac{T_{\text{no-speculation}} - T_{\text{speculation}}}{T_{\text{no-speculation}}} \times 100$$

## 4 Components Supporting Speculative Execution in ARM Processors

### 4.1 Branch Predictors

- Predicts the direction (taken/not-taken) and target of branch instructions.
- Reduces uncertainty in pipeline execution.

### 4.2 Branch Target Buffer (BTB)

- Caches target addresses of recently executed branches.
- Provides quick access to likely branch targets.

### 4.3 Reorder Buffer (ROB)

- Holds speculative results until they are verified as correct.
- Ensures precise exception handling.

### 4.4 Return Address Stack (RAS)

- Tracks return addresses for function calls, ensuring accurate predictions for nested calls.

## 5 ARM Speculative Execution Mechanism

### 5.1 Step-by-Step Process

#### 1. Fetch Instructions Speculatively:

- Fetch instructions along the predicted path.

#### 2. Execute Speculatively:

- Execute the instructions without committing results.
- Maintain speculative results in temporary buffers (e.g., ROB).

#### 3. Verify Prediction:

- Check the branch prediction outcome.
- If correct, commit the results.
- If incorrect, flush the pipeline and re-execute.

## 6 Equations for Speculative Execution in ARM

### 6.1 Correct Prediction

If  $P_c$  is the probability of a correct branch prediction:

$$\text{Execution Time (Correct Prediction)} = I \times C$$

### 6.2 Incorrect Prediction

If  $P_m$  is the probability of a misprediction:

$$\text{Misprediction Time} = P_m \times T_m \times \text{Branch Instructions}$$

### 6.3 Overall Execution Time

Combining both cases:

$$T_{\text{speculative}} = I \times C + P_m \times T_m \times \text{Branch Instructions}$$

## 7 Security Implications of Speculative Execution

### 7.1 Side-Channel Attacks

- Exploit speculative instructions that access restricted data.
- Example: Spectre and Meltdown attacks.

## 7.2 Mitigations in ARM

- **Speculation Barriers:** ARM introduced 'CSDB' (Context Synchronization Barriers) to limit speculative execution.
- **Branch Target Isolation (BTI):** Prevents exploitation of mispredicted branch targets.

## 8 ARM Architectures with Speculative Execution

### 8.1 ARMv7

- Uses basic speculative execution.
- Optimized for mobile devices with short pipelines.

### 8.2 ARMv8

- Advanced dynamic branch predictors.
- Speculative execution supports **AArch64** (64-bit).

### 8.3 ARMv9

- Introduces **neural-based predictors** for better accuracy.
- Mitigates speculative vulnerabilities using **branch protection mechanisms**.

## 9 Advantages of Speculative Execution

1. **Improved Performance:**
  - Keeps the pipeline busy, maximizing throughput.
2. **Reduced Branch Penalty:**
  - Mitigates the performance loss caused by branch instructions.
3. **Efficient Resource Utilization:**
  - Ensures all pipeline stages remain active.

## 10 Challenges of Speculative Execution

1. **Branch Misprediction Penalty:**
  - Requires flushing the pipeline, which wastes cycles.
2. **Hardware Complexity:**

- Speculative execution introduces additional hardware like BTBs, ROBs, and RAS.

### 3. Security Concerns:

- Speculative instructions may leak sensitive data through side channels.

## 11 Example: Speculative Execution in ARM Assembly

```

CMP R1, R2          ; Compare two registers
BEQ TARGET          ; Branch to TARGET if equal
MOV R3, #1           ; Speculatively execute this instruction
TARGET:
ADD R4, R3, #10      ; Add 10 to R3

```

### Explanation:

- The processor speculates whether the branch (BEQ) will be taken.
- Executes MOV R3, #1 speculatively.
- If the prediction is correct, the results are committed. Otherwise, the pipeline is flushed.

## 12 Metrics for Evaluating Speculative Execution

### 12.1 Speculation Efficiency (SE)

$$SE = \frac{\text{Correct Speculative Instructions}}{\text{Total Speculative Instructions}} \times 100$$

### 12.2 Branch Prediction Accuracy (BPA)

$$BPA = \frac{\text{Correct Branch Predictions}}{\text{Total Branches}} \times 100$$

### 12.3 Misprediction Penalty (MP)

$$MP = T_m \times P_m$$

## 13 Future of Speculative Execution in ARM

### 1. Neural Predictors:

- Use machine learning models to improve branch prediction accuracy.

### 2. Context-Aware Execution:

- Adapt speculative execution based on workload characteristics.

### 3. Secure Speculation:

- Develop hardware and software mechanisms to prevent speculative vulnerabilities.

## 14 Conclusion

Speculative execution is a cornerstone of modern ARM processors, enabling high-performance computing by reducing pipeline stalls. Despite its challenges, such as branch misprediction penalties and security risks, ARM's advanced predictors and mitigation techniques ensure efficient and secure execution for a wide range of applications.