

Broadcast-Based Cache Writes, Snoop Bus, and Directory Coherence

1. Introduction

In shared-memory multiprocessor systems, maintaining cache coherence is crucial to ensure that processors have a consistent view of shared data. This requires mechanisms for propagating updates or invalidations to other caches when data is written. The three primary techniques to achieve this are **broadcast-based cache writes**, **snoop buses**, and **directory-based coherence**.

2. Broadcast-Based Cache Writes

2.1 Definition

In broadcast-based cache writes, when a processor writes to a cache line, the operation is broadcasted to all other processors in the system. This ensures all caches are updated or invalidated with minimal delay.

2.2 Types of Broadcast-Based Writes

1. Write-Invalidate Protocol:

- The writing processor invalidates the cache line in all other processors.
- Future reads from other processors result in a cache miss.

2. Write-Update Protocol:

- The new data is broadcasted, and all processors update their caches.

2.3 Mathematical Model

The time required for a broadcast-based cache write can be expressed as:

$$T_{\text{broadcast}} = T_{\text{local}} + T_{\text{propagation}} + T_{\text{update/invalidation}}$$

Where:

- T_{local} : Time to perform the write in the local cache.
- $T_{\text{propagation}}$: Time to propagate the write to other processors.
- $T_{\text{update/invalidation}}$: Time taken by other processors to update or invalidate their caches.

Communication Overhead: The bandwidth $B_{\text{broadcast}}$ required for broadcast:

$$B_{\text{broadcast}} = N \cdot S_{\text{msg}} / T_{\text{cycle}}$$

Where:

- N : Number of processors.
- S_{msg} : Size of the broadcast message.
- T_{cycle} : Time per communication cycle.

3. Snoop Bus

3.1 Definition

A **snoop bus** is a hardware-based mechanism for cache coherence in shared-memory systems. All processors monitor (or "snoop") a shared communication bus for coherence events, such as reads or writes, to maintain a consistent cache state.

3.2 How It Works

1. Shared Bus:

- Processors share a communication bus for memory operations.

2. Monitoring:

- Each cache monitors the bus for read or write operations related to its cached data.

3. Coherence Actions:

- If a processor detects a relevant operation, it updates or invalidates its cache line as needed.

3.3 Coherence Protocols for Snoop Bus

1. MESI Protocol:

- Modified, Exclusive, Shared, Invalid states.
- Ensures coherence through state transitions based on snoop bus events.

2. MOESI Protocol:

- Adds an "Owned" state to MESI for efficient sharing of modified data.

3.4 Mathematical Model

The total time for a snoop bus operation:

$$T_{\text{snoop}} = T_{\text{bus_access}} + T_{\text{cache_action}}$$

Where:

- $T_{\text{bus_access}}$: Time to detect and process the snoop request.
- $T_{\text{cache_action}}$: Time to update or invalidate the cache line.

Bandwidth Usage: The snoop bus bandwidth B_{snoop} :

$$B_{\text{snoop}} = \frac{N_{\text{requests}} \cdot S_{\text{msg}}}{T_{\text{cycle}}}$$

Where:

- N_{requests} : Number of memory requests monitored.
- S_{msg} : Size of the coherence message.

4. Directory-Based Coherence

4.1 Definition

Directory-based coherence uses a centralized or distributed directory to track the state of each cache line in the system. The directory maintains information about which caches hold copies of a memory block and whether the block is modified, shared, or invalid.

4.2 How It Works

1. Directory Entries:

- Each memory block has an associated directory entry.
- Tracks which processors have cached the block and its state (Modified, Shared, Invalid).

2. Coherence Actions:

- Read or write requests are sent to the directory.
- The directory coordinates updates or invalidations across caches.

4.3 Directory States

1. **Uncached:** No processor has the cache block.
2. **Shared:** The block is cached by multiple processors.
3. **Exclusive:** The block is cached by one processor and matches memory.
4. **Modified:** The block is cached by one processor and differs from memory.

4.4 Mathematical Model

Directory Storage Overhead: The directory storage requirement:

$$S_{\text{directory}} = N_{\text{blocks}} \cdot (P + S_{\text{state}})$$

Where:

- N_{blocks} : Number of memory blocks.
- P : Number of processors (bits to identify processors caching a block).
- S_{state} : Storage required for the state of each block.

Total Access Time: The access time for a directory-based coherence system:

$$T_{\text{directory}} = T_{\text{lookup}} + T_{\text{propagation}} + T_{\text{update/invalidation}}$$

Where:

- T_{lookup} : Time to lookup the directory.
- $T_{\text{propagation}}$: Time to propagate updates or invalidations.
- $T_{\text{update/invalidation}}$: Time taken by caches to update or invalidate data.

Scalability: Directory-based systems scale better than broadcast or snoop systems as the number of processors increases, with communication complexity proportional to $O(\log P)$.

5. Comparison of Techniques

Feature	Broadcast-Based Writes	Snoop Bus	Directory Coherence
Scalability	Limited (high bandwidth usage)	Moderate (bus contention)	High (centralized/distributed)
Communication	High (broadcast to all)	Medium (snoop traffic)	Low (targeted messages)
Memory Overhead	Low	Low	High (directory storage)
Latency	Low for small systems	Medium	Low (scalable systems)

6. Conclusion

Broadcast-based writes, snoop buses, and directory-based coherence are three key techniques for maintaining cache coherence in shared-memory multiprocessor systems. While broadcast and snoop-based systems are simpler and effective for small-scale systems, directory-based coherence scales better for large systems due to reduced communication overhead. Mathematical models help optimize these mechanisms by balancing latency, bandwidth, and storage requirements.