

Basic Concepts – Centralized Parallel and Distributed

Algorithms, Examples of distributed systems,

Correctness proof of distributed algorithms; Design

issues in a distributed system; Distributed Computing

Models-Distributed Operating System, Network

Operating System, Middleware based Distributed

System; Client Server model – 2 tier and n-tier CS

model.

Basic Concepts of Distributed Systems and Cloud Computing

1. Centralized, Parallel, and Distributed Systems

1.1 Centralized Systems

- A single computing unit (server) handles all tasks.
- Users access it through terminals.
- Examples: Mainframe computing, early Unix systems.

1.2 Parallel Systems

- Multiple processors work together in a tightly coupled manner to perform a task.
- Shared memory or high-speed interconnects allow communication.
- Used in supercomputers and high-performance computing.

1.3 Distributed Systems

- Independent computers communicate over a network to achieve a common goal.
- No global clock; processes run asynchronously.
- Examples: Google Cloud, Hadoop, Blockchain networks.

2. Examples of Distributed Systems

1. **Google Search Engine** – Uses distributed data centers worldwide to provide search results quickly.
2. **Bitcoin and Blockchain** – A decentralized, peer-to-peer ledger system.
3. **Distributed File Systems** – NFS (Network File System), Google File System (GFS).
4. **Cloud Computing Platforms** – Amazon Web Services (AWS), Microsoft Azure.

3. Correctness Proof of Distributed Algorithms

Distributed algorithms must satisfy:

- **Safety** – The system should not reach an incorrect state.
- **Liveness** – The algorithm should eventually complete its task.
- **Termination** – The algorithm should reach a final state.

Example of Correctness Proof

For **Mutual Exclusion in Distributed Systems**:

1. **Safety**: Ensure that only one process enters the critical section at a time.
2. **Liveness**: Every requesting process gets a chance to enter the critical section.
3. **Fairness**: No process is indefinitely delayed (no starvation).

4. Design Issues in a Distributed System

1. **Transparency**
 - **Access Transparency** – Users should not know if data is remote or local.
 - **Location Transparency** – The physical location of resources is hidden.
 - **Replication Transparency** – Users should not be aware of multiple copies of data.
2. **Fault Tolerance**
 - The system should handle failures without affecting overall functionality.
 - Techniques: Replication, checkpointing, message logging.
3. **Scalability**
 - The system should efficiently handle an increasing number of nodes.
4. **Concurrency**
 - Multiple users should access the system simultaneously without conflicts.
5. **Security**
 - Authentication, Authorization, Encryption, and Secure Communication.

5. Distributed Computing Models

5.1 Distributed Operating System (DOS)

- A single OS manages multiple nodes as one unit.
- Example: Amoeba, Plan 9, Google's Borg OS.

5.2 Network Operating System (NOS)

- Computers operate independently but communicate via a network.
- Example: Microsoft Windows Server, Novell NetWare.

5.3 Middleware-Based Distributed System

- Middleware sits between OS and applications, facilitating communication.
 - Example: CORBA, Java RMI, Apache Kafka.
-

6. Client-Server Model

The client-server model is the backbone of distributed systems.

6.1 2-Tier Model

- Client directly communicates with the server.
- Example: Database applications, where the frontend (client) directly talks to a backend (MySQL, PostgreSQL).

6.2 N-Tier Model

- Involves multiple layers like Presentation, Business Logic, and Data Storage.
 - Example: Web Applications using Angular (frontend), Spring Boot (backend), and MySQL (database).
-

**Shared memory and message passing systems,
synchronous and asynchronous systems, transient and
persistent system, point-to-point and publish-subscribe
systems; Distributed shared memory and memory
consistency models; Group communication - multicast
and anycast communication, message ordering and
message delivery; Remote Procedure Call (RPC).**

Case studies – Socket, MPI, JMS and RMI

Distributed System and Cloud Computing – Advanced Concepts

1. Shared Memory and Message Passing Systems

1.1 Shared Memory Systems

- Multiple processes share a common memory space.
- Processes communicate by reading/writing to shared memory.
- Requires synchronization mechanisms like semaphores and locks.
- **Example:** OpenMP (for parallel programming).

1.2 Message Passing Systems

- Processes communicate by sending and receiving messages.
- No shared memory; suitable for distributed environments.
- **Example:** MPI (Message Passing Interface), TCP/IP sockets.

Comparison:

Feature	Shared Memory	Message Passing
Communication Overhead	Low (direct memory access)	High (data transmission over network)
Synchronization	Complex (requires locks)	Easier (uses message queues)
Scalability	Limited (memory contention)	High (independent processes)

2. Synchronous and Asynchronous Systems

2.1 Synchronous Systems

- Processes execute in lockstep with a global clock.
- **Advantages:** Predictable execution, easier debugging.
- **Disadvantages:** Blocking delays if one process is slow.
- **Example:** Real-time systems, Distributed databases with two-phase commit.

2.2 Asynchronous Systems

- Processes execute independently without a global clock.
 - **Advantages:** More flexible, non-blocking operations.
 - **Disadvantages:** Harder to debug due to race conditions.
 - **Example:** Internet-based systems, email servers.
-

3. Transient and Persistent Systems

3.1 Transient Systems

- Messages exist only while processes are running.
- If a message is not received immediately, it is lost.

- **Example:** UDP-based communication (video streaming).

3.2 Persistent Systems

- Messages are stored until the receiver retrieves them.
- **Example:** Message queues (Kafka, RabbitMQ), Email systems.

4. Point-to-Point and Publish-Subscribe Systems

4.1 Point-to-Point Communication

- Direct communication between sender and receiver.
- Example: TCP sockets, direct RPC calls.

4.2 Publish-Subscribe Systems

- A publisher sends messages without knowing specific receivers (subscribers).
- Subscribers register for specific message topics.
- Example: Kafka, JMS (Java Message Service), MQTT.

Comparison:

Feature	Point-to-Point	Publish-Subscribe
Communication	Direct	Indirect
Scalability	Limited	High
Use Case	Request-Response	Event-driven systems

5. Distributed Shared Memory (DSM) and Memory Consistency Models

5.1 Distributed Shared Memory (DSM)

- Simulates a shared memory system across distributed nodes.
- Data consistency must be maintained.
- Example: NUMA (Non-Uniform Memory Access), TreadMarks (a DSM system).

5.2 Memory Consistency Models

1. **Strict Consistency** – Read always returns the latest write (ideal but impractical).
2. **Sequential Consistency** – All operations appear in the same order to all processes.
3. **Causal Consistency** – Related writes appear in order, but independent operations may be unordered.
4. **Eventual Consistency** – Updates propagate gradually; used in NoSQL databases (Cassandra, DynamoDB).

6. Group Communication

6.1 Multicast Communication

- One-to-many communication.
- Used for group messaging (e.g., video conferencing).
- Example: IP multicast, PGM (Pragmatic General Multicast).

6.2 Anycast Communication

- Sender sends data to the nearest available receiver.
- Example: CDN (Content Delivery Networks) where users get content from the nearest server.

7. Message Ordering and Message Delivery

7.1 Message Ordering

1. **FIFO Ordering** – Messages from a sender are received in the order sent.
2. **Causal Ordering** – Messages related to each other follow a specific order.
3. **Total Ordering** – All messages appear in the same order to all recipients.

7.2 Message Delivery

- **Reliable Delivery:** Ensures messages are not lost.
- **Atomic Delivery:** Messages are either delivered to all or none.

8. Remote Procedure Call (RPC)

- Allows a process to execute code on a remote server as if it were a local function.
- **Steps in RPC:**
 1. Client calls a local stub function.
 2. Stub function marshals (packs) data and sends a request to the server.
 3. Server receives the request, executes the function, and sends back the result.
 4. Client stub receives the response and returns the result.
- **Example Implementations:** gRPC, Java RMI, CORBA.

9. Case Studies

9.1 Socket Programming

- Provides low-level network communication using TCP/UDP.
- Used for building custom communication protocols.
- Example: Python `socket` module, Java `Socket` class.

9.2 MPI (Message Passing Interface)

- Used for high-performance parallel computing.
- Provides primitives for sending/receiving messages between processes.
- Example: Used in scientific computing applications.

9.3 JMS (Java Message Service)

- A publish-subscribe model for Java applications.
- Supports durable message storage.
- Example: Apache ActiveMQ, RabbitMQ.

9.4 RMI (Remote Method Invocation)

- Java's built-in RPC framework for distributed objects.
 - Supports method calls over a network.
 - Example: A Java server exposes a method, and a remote client calls it.
-

Clock, events and process state in distributed systems;

Physical clock synchronization, Logical clocks and

Vector clocks; Distributed Snapshots – Global States;

Distributed mutual exclusions; Leader election

algorithms; Distributed transactions; Deadlocks in

distributed systems;

Distributed Systems: Clocks, Synchronization, and Coordination

1. Clock, Events, and Process State in Distributed Systems

1.1 Clocks in Distributed Systems

- Each node in a distributed system has its own clock.

- Clocks may drift due to different clock speeds, leading to time inconsistencies.

1.2 Events in Distributed Systems

- **Local Events:** Actions occurring within a single process (e.g., variable update).
- **Message Passing Events:** Sending and receiving messages between processes.
- **Global Events:** Events affecting multiple processes (e.g., system failure).

1.3 Process State

- Each process has its own execution state.
 - The state consists of local variables, pending messages, and the logical clock value.
-

2. Physical Clock Synchronization

Clock synchronization ensures consistency across nodes.

2.1 Christian's Algorithm (Centralized Approach)

- A time server provides clock values to all nodes.
- Each node adjusts its time using:

$$T = T_{\text{server}} + \frac{(T_{\text{request sent}} - T_{\text{response received}})}{2}$$

- Drawback: Single point of failure (if the server crashes).

2.2 Berkeley's Algorithm (Average Consensus Approach)

- The master node asks all nodes for their time.
- It computes the average and tells all nodes to adjust accordingly.
- Works well for LANs but not in large-scale distributed systems.

2.3 Network Time Protocol (NTP) (Hierarchy-Based Approach)

- A hierarchical system where stratum-1 servers sync with atomic clocks.
 - Uses offset and round-trip delay calculations for precision.
 - Widely used for internet-scale clock synchronization.
-

3. Logical Clocks and Vector Clocks

3.1 Logical Clocks (Lamport's Timestamps)

- Instead of real-time, processes use logical timestamps.
- **Rules:**
 1. Each process maintains a counter (LC).
 2. Before an event, increment LC .
 3. On sending a message, send LC .
 4. On receiving a message, update $LC = \max(LC, received_LC) + 1$.
- Ensures **causal ordering** but does not detect concurrency.

3.2 Vector Clocks

- Each process maintains a vector $[P_1, P_2, \dots, P_n]$.
 - When a process sends a message, it updates its own entry.
 - On receiving, it updates by taking the **maximum of corresponding values**.
 - **Advantage:** Detects concurrency (unlike Lamport's Clocks).
 - **Drawback:** Space complexity $O(n)$.
-

4. Distributed Snapshots – Global State

- Captures a consistent global state of a distributed system.
 - **Chandy-Lamport Algorithm:**
 1. A process initiates the snapshot and records its state.
 2. It sends a marker to all other processes.
 3. When a process receives a marker for the first time, it records its state and forwards the marker.
 4. The process records incoming messages until all markers are received.
 - **Use Case:** Checkpointing in distributed databases.
-

5. Distributed Mutual Exclusion

Mutual exclusion ensures only one process accesses a critical section (CS) at a time.

5.1 Ricart-Agrawala Algorithm (Timestamp-Based Approach)

1. A process sends a request message with a timestamp.
 2. Other processes reply with **OK** if they have a later timestamp or are not in CS.
 3. Process enters CS after receiving **OK** from all.
 4. On exit, it sends a **release** message.
- **Advantage:** Message complexity $O(N)$.
 - **Drawback:** High overhead in large systems.

5.2 Token-Based Algorithms

- A **unique token** grants access to the CS.

- Token circulates among processes.
 - Example: **Token Ring Algorithm** (a logical ring structure).
-

6. Leader Election Algorithms

Leader election is essential for coordination in distributed systems.

6.1 Bully Algorithm

1. A process detects a leader failure and initiates an election.
 2. It sends an **Election** message to higher-ID processes.
 3. If no response, it declares itself the leader.
 4. If a higher-ID process responds, it takes over.
- **Drawback:** High message overhead.

6.2 Ring Algorithm

1. Processes are arranged in a logical ring.
 2. A process initiates an election by passing an election message.
 3. The highest-ID process is elected as the leader.
- **Advantage:** Less message complexity than Bully Algorithm.
-

7. Distributed Transactions

A transaction ensures atomicity in a distributed system.

7.1 ACID Properties in Distributed Transactions

- **Atomicity:** Either all operations commit, or none do.
- **Consistency:** The system moves from one valid state to another.
- **Isolation:** Transactions do not interfere.
- **Durability:** Once committed, changes persist.

7.2 Two-Phase Commit (2PC)

1. **Prepare Phase:** Coordinator asks all nodes if they can commit.
 2. **Commit Phase:** If all nodes agree, commit the transaction; otherwise, abort.
- **Issue:** If the coordinator crashes, the system may hang.

7.3 Three-Phase Commit (3PC)

- Adds a **pre-commit phase** before commit.
 - Ensures that if a crash occurs, nodes can safely recover.
-

8. Deadlocks in Distributed Systems

A **deadlock** occurs when processes wait indefinitely for resources held by others.

8.1 Conditions for Deadlock (Coffman's Conditions)

1. **Mutual Exclusion** – Resources are non-shareable.
2. **Hold and Wait** – Processes hold some resources while waiting for others.
3. **No Preemption** – Resources cannot be forcibly taken.
4. **Circular Wait** – A circular chain of waiting exists.

8.2 Deadlock Detection Algorithm

- Construct a **Wait-for Graph (WFG)** where nodes represent processes and edges represent dependencies.
- A cycle in the WFG means a deadlock exists.

8.3 Deadlock Prevention Techniques

- **Timeout-Based**: Abort transactions if they exceed a threshold.
 - **Request All at Once**: Processes request all required resources upfront.
 - **Resource Ordering**: Impose an order to prevent circular waits.
-

Fault Tolerance

Basic concepts, fault models, consensus problems and

its applications, commit protocols, voting protocols,

Checkpointing and recovery, reliable communication

Fault Tolerance in Distributed Systems

1. Basic Concepts of Fault Tolerance

Fault tolerance is the ability of a system to continue operating correctly despite failures. It ensures system reliability, availability, and maintainability.

1.1 Types of Faults

1. **Transient Faults:** Temporary errors that disappear after some time (e.g., network glitches).
2. **Intermittent Faults:** Occur unpredictably due to unstable components (e.g., loose connections).
3. **Permanent Faults:** Hardware or software failures requiring repair or replacement (e.g., disk failure).

1.2 Fault Tolerance Techniques

- **Redundancy:** Extra components to take over in case of failure.
 - **Replication:** Multiple copies of data or processes running concurrently.
 - **Checkpointing:** Saving system state periodically for recovery.
 - **Recovery Mechanisms:** Rollback to a consistent state after a failure.
-

2. Fault Models

Fault models classify failures based on their behavior and impact.

2.1 Crash Fault Model

- A process **fails by stopping** and does not recover.
- Other processes must detect and handle the failure.
- Example: Node failures in cloud computing.

2.2 Omission Fault Model

- A process fails to send or receive messages.
- Example: Network packet loss in unreliable channels.

2.3 Byzantine Fault Model

- A process behaves arbitrarily, possibly sending conflicting messages.
- **Example:** Malicious attacks, software bugs causing incorrect computations.
- **Solution:** Byzantine Fault Tolerant (BFT) algorithms like PBFT (Practical Byzantine Fault Tolerance).

2.4 Timing Fault Model

- A process does not meet timing constraints.
 - Example: Real-time systems where delayed responses lead to failures.
-

3. Consensus Problems and Applications

Consensus ensures that all non-faulty processes agree on a single decision, even in the presence of failures.

3.1 Requirements for Consensus

1. **Termination** – Every correct process eventually decides on a value.
2. **Agreement** – All correct processes agree on the same value.
3. **Integrity** – The agreed value must have been proposed by a process.
4. **Validity** – The decision must be based on actual inputs.

3.2 FLP Impossibility Theorem

- In an **asynchronous** system, achieving consensus is impossible if even **one process can fail**.
- **Workaround:** Use probabilistic or partially synchronous models (e.g., Paxos, Raft).

3.3 Consensus Algorithms

- **Paxos:** A widely used consensus algorithm in distributed systems.
- **Raft:** Simplifies Paxos by using a leader-based approach.
- **Byzantine Agreement:** Extends consensus to Byzantine faults (e.g., PBFT).

3.4 Applications of Consensus

- Distributed databases (e.g., Google Spanner).
 - Blockchain (e.g., Bitcoin uses Proof-of-Work for consensus).
 - Leader election in distributed systems.
-

4. Commit Protocols

Commit protocols ensure atomic transactions in distributed databases.

4.1 Two-Phase Commit (2PC)

- Ensures that either all nodes commit or none do.

Steps in 2PC:

1. **Prepare Phase:** Coordinator asks participants if they can commit.
2. **Commit Phase:** If all agree, coordinator sends a commit; otherwise, it aborts.

Drawbacks of 2PC:

- If the coordinator crashes, the system may hang.
- Requires strict synchronization, leading to high latency.

4.2 Three-Phase Commit (3PC)

- Adds a **pre-commit phase** to reduce blocking issues.

Steps in 3PC:

1. **Prepare Phase** – Coordinator asks nodes if they are ready.
2. **Pre-Commit Phase** – If all agree, coordinator sends a pre-commit message.
3. **Commit Phase** – If no failures occur, the final commit is sent.

Advantage of 3PC:

- **Non-blocking:** Nodes can safely recover from coordinator failure.
-

5. Voting Protocols

Voting protocols ensure fault tolerance by requiring agreement from a majority of processes.

5.1 Majority Voting Protocol

- A transaction is committed if more than 50% of nodes approve.
- Used in **quorum-based systems** (e.g., ZooKeeper).

5.2 Weighted Voting Protocol

- Different nodes have different weights in the voting process.
 - Useful in **heterogeneous distributed systems**.
-

6. Checkpointing and Recovery

Checkpointing saves the system state periodically to enable recovery from failures.

6.1 Types of Checkpointing

1. **Coordinated Checkpointing** – All processes take a checkpoint simultaneously.
 - Avoids inconsistency but requires synchronization.
2. **Uncoordinated Checkpointing** – Processes checkpoint independently.
 - Can lead to **domino effect**, where old checkpoints must be rolled back.
3. **Incremental Checkpointing** – Saves only changes since the last checkpoint.
 - Reduces storage overhead.

6.2 Recovery Techniques

- **Rollback-Recovery:** Restores a process to its last checkpoint.
- **Message Logging:** Logs messages for replaying after a failure.

7. Reliable Communication in Distributed Systems

Reliable communication ensures that messages are delivered correctly and in order.

7.1 Challenges in Reliable Communication

- Network failures (packet loss, duplication, or corruption).
- Process failures (crashes, Byzantine behavior).

7.2 Techniques for Reliable Communication

1. **Acknowledgment-based Protocols**
 - **Stop-and-Wait:** Sender waits for an acknowledgment (ACK) before sending the next message.
 - **Sliding Window Protocol:** Allows multiple messages in transit (e.g., TCP).
2. **Retransmission Mechanisms**
 - **Automatic Repeat reQuest (ARQ):** Retransmits lost packets.
 - **Forward Error Correction (FEC):** Adds redundancy for error detection.
3. **Ordered Delivery**
 - **FIFO Ordering:** Messages from a sender are received in order.
 - **Causal Ordering:** Messages maintain causality constraints.
 - **Total Ordering:** All recipients receive messages in the same order.

Summary Table

Concept	Key Points	Example Algorithms
Fault Models	Crash, Omission, Byzantine, Timing	PBFT for Byzantine faults
Consensus	Agreement, Integrity, FLP Impossibility	Paxos, Raft, Byzantine Agreement
Commit Protocols	Ensure atomic transactions	2PC, 3PC
Voting Protocols	Majority-based decision making	Quorum-based voting
Checkpointing	Saves system state	Coordinated, Uncoordinated, Incremental
Reliable Communication	Ensures correct and ordered delivery	TCP, ARQ, Sliding Window

Big Data; Relation between Big Data and Distributed

File System; Map Reduce Programming Paradigm.

Case Studies: HDFS and Hadoop Map Reduce

Distributed File System (DFS) and Big Data

1. Distributed File System (DFS)

A **Distributed File System (DFS)** is a system that allows files to be stored across multiple machines while appearing as a single logical file system to users. DFS provides:

- **Scalability:** Handles large amounts of data.
- **Fault tolerance:** Ensures availability despite node failures.
- **Concurrency:** Supports multiple users accessing files simultaneously.

1.1 Characteristics of DFS

1. **Transparency**
 - **Access Transparency:** Users access remote files like local files.
 - **Location Transparency:** File location is hidden from users.
 - **Replication Transparency:** Multiple copies of a file exist without user intervention.
 - **Failure Transparency:** System recovers automatically from failures.
 2. **Scalability**
 - Handles a growing amount of data and users.
 3. **Fault Tolerance**
 - Uses replication, checkpoints, and recovery techniques to handle failures.
 4. **Concurrency Control**
 - Ensures consistency when multiple users access or modify files.
-

2. Big Data and Its Relation to DFS

2.1 What is Big Data?

Big Data refers to massive datasets that are difficult to store, process, and analyze using traditional methods. It is characterized by:

- **Volume:** Large-scale data storage.
- **Velocity:** High-speed data processing.
- **Variety:** Different formats (structured, semi-structured, unstructured).
- **Veracity:** Data accuracy and reliability.
- **Value:** Extracting useful insights from data.

2.2 Why DFS for Big Data?

Big Data requires a **scalable and distributed** storage system. DFS efficiently handles:

- **Storage of massive datasets.**
- **Efficient data retrieval** across multiple nodes.
- **Fault tolerance** through replication.
- **Parallel processing support** (used in **MapReduce**).

Example: Hadoop Distributed File System (HDFS) is widely used for handling Big Data.

3. MapReduce Programming Paradigm

MapReduce is a programming model for processing large datasets in parallel across a distributed system.

3.1 Key Concepts

- **Map Phase:** Divides input data into smaller chunks and processes them in parallel.
- **Shuffle Phase:** Groups and sorts intermediate results.
- **Reduce Phase:** Aggregates and processes the intermediate data.

3.2 Example Workflow

Problem: Count word occurrences in a large dataset.

1. **Map Phase:**
 - Input: ["Hello world", "Hello Hadoop"]
 - Output: (Hello, 1), (world, 1), (Hello, 1), (Hadoop, 1)
 2. **Shuffle Phase:**
 - Groups words together:
(Hello, [1,1]), (world, [1]), (Hadoop, [1])
 3. **Reduce Phase:**
 - Aggregates counts:
(Hello, 2), (world, 1), (Hadoop, 1)
-

4. Case Studies: HDFS and Hadoop MapReduce

4.1 Hadoop Distributed File System (HDFS)

HDFS is a DFS designed for handling Big Data.

Key Features

- **Master-Slave Architecture:**
 - **NameNode** (Master) manages metadata and file structure.
 - **DataNodes** (Slaves) store actual file blocks.

- **Block Storage:** Files are divided into **blocks (128 MB default)** and stored across multiple nodes.
- **Replication:** Default **3x replication** for fault tolerance.
- **Write-Once-Read-Many (WORM) Model:** Optimized for large sequential reads.

HDFS Architecture

Client → NameNode → DataNodes (Stores Blocks)

- Clients interact with **NameNode** to locate files.
- Data is stored and replicated across **DataNodes**.

4.2 Hadoop MapReduce

Hadoop MapReduce uses a DFS (like HDFS) for parallel processing of data.

Advantages of Hadoop MapReduce

- **Parallel Execution:** Tasks are executed concurrently on multiple nodes.
- **Fault Tolerant:** Failed tasks are automatically restarted.
- **Scalable:** Works on clusters of thousands of machines.

5. Summary Table

Concept	Key Points	Example
DFS	Stores files across multiple nodes	HDFS
Big Data	Large-scale, high-speed data processing	Apache Spark, Hadoop
MapReduce	Distributed processing model	Word count, Log analysis
HDFS	Distributed file storage, Replication	Used in Hadoop
Hadoop MapReduce	Processes data in parallel	Log analysis, Data mining

Cloud Computing

Definition of Cloud Computing; Cloud Service Models;

Cloud Deployment Models; Virtualization Technologies.

Case Studies: AWS, GCP and Windows Azure

Cloud Computing

1. Definition of Cloud Computing

Cloud Computing is the **on-demand** delivery of computing resources such as servers, storage, databases, networking, software, and analytics over the **internet**. It allows users to access and manage resources **without maintaining physical infrastructure**.

1.1 Characteristics of Cloud Computing

1. **On-Demand Self-Service** – Users can provision resources as needed.
 2. **Broad Network Access** – Services are accessible from anywhere.
 3. **Resource Pooling** – Multiple users share computing resources.
 4. **Rapid Elasticity** – Resources can scale up/down automatically.
 5. **Measured Service** – Users pay only for what they use (pay-as-you-go).
-

2. Cloud Service Models (SPI Model)

Cloud services are categorized into **three main models**, also called the **SPI model**:

Service Model	Description	Examples
Software as a Service (SaaS)	Provides software applications over the internet. Users do not manage the infrastructure.	Gmail, Google Drive, Dropbox
Platform as a Service (PaaS)	Provides a platform for developers to build applications without managing underlying hardware.	Google App Engine, AWS Elastic Beanstalk
Infrastructure as a Service (IaaS)	Provides virtualized computing resources like servers, storage, and networking.	AWS EC2, Google Compute Engine

2.1 Comparison of Service Models

- **SaaS**: End-users consume applications via the cloud.
 - **PaaS**: Developers deploy and manage applications.
 - **IaaS**: IT administrators manage servers, storage, and networking.
-

3. Cloud Deployment Models

Deployment Model	Description	Use Case
Public Cloud	Resources are owned and managed by a third-party provider. Shared among multiple users.	Startups, Web applications
Private Cloud	Dedicated infrastructure for a single organization. More control and security.	Banking, Healthcare
Hybrid Cloud	Combination of public and private clouds. Used for sensitive and scalable workloads.	Enterprises, Large-scale businesses

Deployment Model	Description	Use Case
Community Cloud	Shared cloud infrastructure for a specific group of organizations.	Government, Research institutions

3.1 Public vs Private Cloud

Feature	Public Cloud	Private Cloud
Security	Lower	Higher
Cost	Pay-per-use (cheaper)	Expensive
Scalability	High	Limited
Control	Less control	Full control

4. Virtualization Technologies

Virtualization is a key technology that enables cloud computing by creating multiple **virtual instances** of hardware or software.

4.1 Types of Virtualization

- Hardware Virtualization**
 - Uses a **hypervisor** to create Virtual Machines (VMs).
 - Examples:** VMware, VirtualBox, KVM.
- Server Virtualization**
 - Multiple servers run as independent VMs on one physical server.
 - Example:** AWS EC2 instances.
- Storage Virtualization**
 - Abstracts multiple storage devices into a single logical unit.
 - Example:** Amazon S3, Google Cloud Storage.
- Network Virtualization**
 - Virtualizes network resources for better efficiency.
 - Example:** Software-Defined Networking (SDN).
- Desktop Virtualization**
 - Allows remote access to a desktop environment.
 - Example:** Windows Virtual Desktop.

4.2 Hypervisors

A **hypervisor** is software that enables virtualization by managing VMs.

Type	Description	Example
Type-1 (Bare Metal)	Runs directly on hardware	VMware ESXi, Microsoft Hyper-V
Type-2 (Hosted)	Runs on top of an OS	VirtualBox, VMware Workstation

5. Case Studies: AWS, GCP, and Microsoft Azure

5.1 Amazon Web Services (AWS)

AWS is the **largest cloud provider**, offering **IaaS, PaaS, and SaaS** solutions.

- **Compute Services:** EC2 (VMs), Lambda (serverless computing).
- **Storage:** S3 (object storage), EBS (block storage).
- **Database:** RDS (relational DB), DynamoDB (NoSQL).
- **Networking:** VPC, CloudFront (CDN).

Use Case: Netflix uses AWS for streaming content globally.

5.2 Google Cloud Platform (GCP)

GCP provides cloud infrastructure, AI, and analytics tools.

- **Compute:** Compute Engine (VMs), Kubernetes Engine.
- **Storage:** Cloud Storage, Persistent Disk.
- **AI & ML:** TensorFlow, Vertex AI.
- **Big Data:** BigQuery, Dataflow.

Use Case: Spotify uses GCP for machine learning and recommendations.

5.3 Microsoft Azure

Azure is widely used for enterprise solutions and integrates with Microsoft products.

- **Compute:** Azure Virtual Machines, Functions (serverless).
- **Storage:** Blob Storage, Azure Files.
- **AI & Analytics:** Azure Machine Learning, Power BI.
- **Hybrid Cloud:** Azure Stack, Hybrid Networking.

Use Case: BMW uses Azure for IoT and AI-based car analytics.

6. Summary Table

Concept	Key Points	Examples
Cloud Computing	Internet-based computing, scalable & cost-effective	AWS, GCP, Azure
Cloud Service Models	SaaS, PaaS, IaaS	Google Drive (SaaS), AWS EC2 (IaaS)

Concept	Key Points	Examples
Deployment Models	Public, Private, Hybrid, Community	AWS (Public), Azure Stack (Hybrid)
Virtualization	Creates virtual resources	Hypervisors: VMware, KVM
AWS	Largest cloud provider	EC2, S3, Lambda
GCP	Strong AI & ML services	Compute Engine, BigQuery
Azure	Best for enterprise solutions	Azure VM, Blob Storage
<hr/>		
<hr/>		